

Deploy your LLM

Ignacio Heredia

Instituto de Física de Cantabria (CSIC-UC)

Master Data Science (UIMP-UC)

March 2025



iheredia@ifca.unican.es



[IgnacioHeredia](#)

- Motivation
- Docker mini-tutorial
- Deploy a chatbot with ollama
- Deploy a vision model with ollama
- Use ollama from Python
- Build a Retrieval Augmented Generation workflow

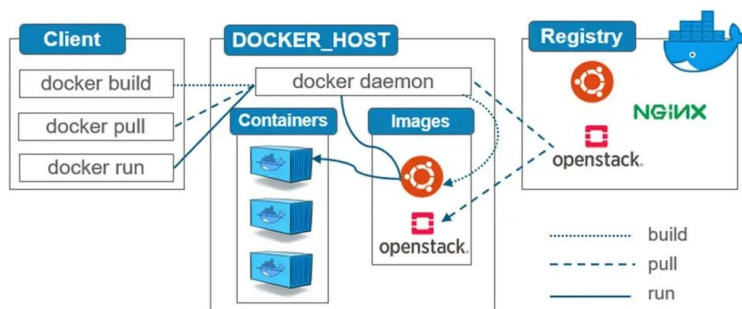
Why deploying models locally?

- 🔍 **privacy**: deploying locally allows you to use LLMs without sending your personal data to a private company that might sell it,
- 👑 **sovereignty**: you won't rely on an external private company.
- 💰 **cost**: using local models is free!
- ⚡ **speed**: when using models without RAG, small local models reply faster than an large LLM hosted somewhere else
- 🌱 **energy**: using small models (usually) is more energy efficient

The main disadvantage is that small local model are less powerful than big ones. But for simple things, they do work surprisingly well.



Docker is an open-source platform that allows developers to build, deploy, run, and manage containerized applications. It uses OS-level virtualization to package software and its dependencies into standardized units called **containers**.



Docker allows developers to package their software and its dependencies in a single unit that can run anywhere.

The most popular registry is [DockerHub](https://hub.docker.com/).

- Docker Engine has [installation guides for Linux](#).
- For Windows/Mac user, you can use it through [Docker Desktop](#).
- Installation steps for [Ubuntu users](#):



Go ahead and install it!

Install using the `apt` repository

Before you install Docker Engine for the first time on a new host machine, you need to set up the Docker `apt` repository. Afterward, you can install and update Docker from the repository.

1. Set up Docker's `apt` repository.

```
# Add Docker's official GPG key:
sudo apt-get update
sudo apt-get install ca-certificates curl
sudo install -m 0755 -d /etc/apt/keyrings
sudo curl -fsSL https://download.docker.com/linux/ubuntu/gpg -o /etc/apt/keyrings/docker.asc
sudo chmod a+r /etc/apt/keyrings/docker.asc

# Add the repository to Apt sources:
echo \
  "deb [arch=$(dpkg --print-architecture) signed-by=/etc/apt/keyrings/docker.asc] https://download.docker.com/linux/ubuntu $(. /etc/os-release && echo "${UBUNTU_CODENAME:-$VERSION_CODENAME}") stable" | \
  sudo tee /etc/apt/sources.list.d/docker.list > /dev/null
sudo apt-get update
```

2. Install the Docker packages.

Latest Specific version

To install the latest version, run:

```
$ sudo apt-get install docker-ce docker-ce-cli containerd.io docker-buildx-plugin docker-compose-plugin
```

3. Verify that the installation is successful by running the `hello-world` image:


```
$ sudo docker run hello-world
```

This command downloads a test image and runs it in a container. When the container runs, it prints a confirmation message and exits.

You have now successfully installed and started Docker Engine.

- **docker pull** <image>
 - **docker images**
 - **docker run** <image>
 - **docker ps**
 - **docker exec -it** <container> <command>
 - **docker rm** <container>
- pull an image from a registry (eg. Dockerhub)
 - list your pulled images
 - create a new container from an image
 - list your deployed containers
 - run a command in a container
 - delete a container

Deploy Llama 3.2 with Ollama

- We are going to deploy an LLM locally with [ollama](#)
- We can choose any model from the [ollama Marketplace](#).
- You should have at least:
 - 8 GB of RAM available to run the 7B models,
 - 16 GB to run the 13B models,
 - 32 GB to run the 33B models.
- For a better latency experience, we are going to deploy the smaller models.
- Let's deploy [Llama 3.2](#) (1B) from  Meta

```
$ docker run -d -v ollama:/root/.ollama -p 11434:11434 --name ollama ollama/ollama
$ docker exec -it ollama /bin/bash
> ollama pull llama3.2:1b
> ollama list
> ollama run llama3.2:1b
> ollama stop llama3.2:1b
```

Deploy DeepSeek-R1 with Ollama

- Let's get fancier and deploy [DeepSeek-R1 7B](#) (Qwen distill)  deepseek

```
> ollama run deepseek-r1:7b
```

- You can see that it is iteratively reflecting over its own responses

```
>>> You're so cool
```

```
<think>
```

```
Alright, the user just said "You're so cool." That's pretty straightforward.  
I need to respond in a friendly way. Maybe acknowledge their comment.
```

```
I should keep it simple and positive without overcomplicating things.
```

```
</think>
```

```
Thank you! I'm glad you like it. How can I assist you further? 😊
```

(⚠) People with low resources can use **deepseek-r1:1.5b**

- We can also deploy a small vision model like IBM's [granite3.2-vision:2b](#) to test the capabilities to analyse images:

```
$ docker run -d -v ollama:/root/.ollama -p 11434:11434 --name ollama ollama/ollama
$ docker exec -it ollama /bin/bash
> apt update && apt-get install wget
> wget -O demo.jpg <image-url>
> ollama run granite3.2-vision:2b
>>> Describe the image in /demo.jpg
```


⚠ Take into account that Vision models can be **really slow**, especially when running on large images using CPU only.

Deploy a Vision model



The image depicts a lion walking across a sandy terrain. The lion is positioned towards the left side of the frame, with its body oriented slightly to the right. It has a robust and muscular build, typical of a male lion. Its mane is long and thick, [...]

Correctly detects lion, as well as its surroundings



The image is a bar chart that displays data on the "Number of **businesses** manufactured" over different months of the year. The x-axis represents the months of the year, labeled as January, February, March, April, [...]

Confuses bus with businesses, rest is pretty fine



The humor in this image stems from the juxtaposition of a young child's innocent **curiosity** with the chaotic backdrop of a fire. The child's expression, which seems to be one of **wonder or mild concern**, contrasts sharply with the [...]

It's not able to pick up complex human expression

The model is able to perform (reasonably) well in a wide variety of tasks. We say the models are **zero-shot**.


- You can use the [ollama Python library](#), to query your model from Python.
- For more advanced workflows, like Retrieval Augmented Generation, you can use the [Llama Index](#) library.

```
$ python -m venv --system-site-packages llama-index-env
$ source llama-index-env/bin/activate
$ pip install ollama
$ pip install llama-index
$ pip install llama-index-llms-ollama
$ deactivate
```

- For example, a simple usecase would be to use an LLM to *fix our data during preprocessing*.
- The processing is easy enough so we can go with a non-thinking model that usually performs faster. In addition, distilled deepseek model are not trained with RL so they can output weird data sometimes.
- Let's deploy [Llama 3.2](#) (3B).

```
$ docker run -d -v ollama:/root/.ollama -p 11434:11434 --name ollama ollama/ollama
$ docker exec -it ollama /bin/bash
> ollama pull llama3.2:3b
> ollama run llama3.2:3b
```



- Using the ollama module:



```
import ollama

faulty_data = "Ths si som slopppy dta tat willl hopefly gt fxd by sart lage languje mdl"
response = ollama.chat(model='llama3.2:3b', messages=[
    {
        "role": "user",
        "content": f"""
The following phrase has mistakes in it, please fix it.
Only reply with the revised text, nothing else:

{faulty_data}
"""
    },
])
print(response['message']['content'])
# This is some sloppy data that will hopefully get fixed by smart large language model.
```

- We can also use it to perform Retrieval Augmented Generation in your personal Knowledge Base (for example, create a  **data** folder containing  [Guía docente - Machine Learning II](#)):





```
from pathlib import Path


from llama_index.core import VectorStoreIndex, SimpleDirectoryReader, Settings
from llama_index.embeddings.huggingface import HuggingFaceEmbedding
from llama_index.llms.ollama import Ollama

Settings.embed_model = HuggingFaceEmbedding(model_name="BAAI/bge-base-en-v1.5")
Settings.llm = Ollama(model="llama3.2:1b", request_timeout=360.0)

data_folder = Path(__file__).parent / "data"
documents = SimpleDirectoryReader(data_folder).load_data()
index = VectorStoreIndex.from_documents(documents)
query_engine = index.as_query_engine()

query = "Who are the lecturers in Machine Learning?"
response = query_engine.query(query)
print(response)
```

- We can also use it to perform Retrieval Augmented Generation in your personal Knowledge Base (for example, create a  **data** folder containing  [Guía docente - Machine Learning II](#)):

The lecturers listed for the Machine Learning course 249 - AUTOMATIC LEARNING II  at the University Master's Degree in Data Science are:

1. Luis Ignacio Santamaría Caballero
2. Steven Johan Maria van Vaerenberg
3. **S**oto Herrera Garcia
4. Javier**a** Diez Sierra
5. Maia**LEN** Iturbide Martinez De Alben**ez**

(*) Correct answers

LUIS IGNACIO SANTAMARIA CABALLERO
STEVEN JOHAN MARIA VAN VAERENBERGH
SIXTO HERRERA GARCIA
JAVIER DIEZ SIERRA
MAIALEN ITURBIDE MARTINEZ DE ALBENIZ

- Errors could come both from PDF parsing and from model size.
- You could also tweak the previous script to change the temperature (ie. make the LLM more creative), to see the citations the model uses as reference, etc.

- LLMs are very easy to deploy locally, thus benefiting from additional privacy,
- Response latency for smaller models is pretty good, at least for pure-text models,
- Using +7B models locally tends to consume too many resources in standard workstations,
- Use cases:
 - Chat works pretty well, allowing to show the thought process
 - Vision models work fine in simple usecases,
 - Effective for “simple” data wrangling that would otherwise be difficult with standard Python modules,
 - Using RAG delivers mixed results depending on the database size,

Questions

Ignacio Heredia

Instituto de Física de Cantabria (CSIC-UC)



iheredia@ifca.unican.es



[IgnacioHeredia](https://github.com/IgnacioHeredia)



Image sources